

國立金門大學

資訊工程學系碩士班

碩士論文

基於對比度受限自適應直方圖均衡化影像增強深度學習與醫

視智能雲端部署之即時乳癌診斷系統

Real-Time Breast Cancer Diagnosis System Using Clahe-
Enhanced Deep Learning and Cloud Deployment via
MedVision-AI

NQU

研究生：Shi-Han Huang (黃世漢)

指導教授：李錫捷 博士

中華民國一五五年六月

國立金門大學碩士學位論文考試審定書

本校理工學院資訊工程學系碩士班

研究生 Shi-Han Huang (黃世漢) 所提之論文

Real-Time Breast Cancer Diagnosis System Using Clahe-Enhanced Deep
Learning and Cloud Deployment via MedVision-AI

業經本委員會評審認可，合於碩士資格水準。

學位考試委員

召 集 人 _____ 簽章

委 員 _____

指 導 教 授 _____

系 所 主 管 _____

中 華 民 國 115 年 6 月 21 日

基於對比度受限自適應直方圖均衡化影像增強深度學習與醫視智能雲

端部署之即時乳癌診斷系統

研究生：SHI-HAN HUANG (黃世漢)

指導教授：李錫捷 博士

國立金門大學資訊工程學系碩士班

摘要

乳腺癌仍然是全球女性死亡的主要原因之一，這凸顯了對高精度、自動化診斷框架的迫切需求。本研究對乳腺惡性腫瘤的分類進行了全面調查，採用了一種結合傳統機器學習（以 K-Nearest Neighbors, K-NN 為代表）與先進深度學習架構（包括 VGG16、VGG19 和 ResNet50）的整合方法。利用包含超過 10,430 個樣本的大型乳房攝影 (DDSM/Mendeley) 數據集，本研究建立了診斷準確性和臨床解釋性的基準。核心創新在於採用 限制對比度自適應直方圖均衡化 (CLAHE) 的多階段預處理流程，該流程專門針對顯現緻密腺體組織中的微小鈣化點進行了優化，這是傳統乳房攝影篩查中的共同挑戰。實驗結果顯示，雖然優化的 K-NN 模型達到了 96% 的穩定準確度，但深度學習集成模型顯著優於傳統方法，其中 CNN + CLAHE 流程達到了 97.84% 的峰值準確度。為了應對醫療 AI 中系統性的「黑盒子」挑戰，本研究整合了 梯度加權類激活映射 (Grad-CAM)，為放射科醫師提供每次分類的精確視覺「證據」。此外，研究成果成功從 Google Colab 的實驗室規模過渡到基於 Google Antigravity IDE 的生產級雲端生態系統。最終部署的 MedVision-AI 平台包含一個客製化的 檢索增強生成 (RAG) 引擎，該引擎將模型置信度與既定的臨床指南相結合，生成自動化病理報告。這些結果證明，雲端整合、可解釋的 AI 框架可以顯著彌合現代診斷中的可及性差距，為城市醫院和資源受限環境中的臨床醫生提供可靠且可擴展的支持。

關鍵詞：乳腺癌診斷；深度學習；限制對比度自適應直方圖均衡化 (CLAHE)；可解釋人工智慧 (XAI)；Grad-CAM；檢索增強生成 (RAG)；MedVision-AI；Google Antigravity；VGG16；ResNet50。

Real-Time Breast Cancer Diagnosis System Using Clahe-Enhanced Deep Learning and Cloud Deployment via MedVision-AI

Student: SHI-HAN HUANG (黃世漢)

Advisor: HIS-CHIEH LEE, PhD.

Master's Degree Program of Computer Science and Information Engineering,
National Quemoy University

ABSTRACT

Breast cancer continues to be a primary contributor to female mortality worldwide, highlighting the critical imperative for high-precision, automated diagnostic frameworks. This research presents a comprehensive investigation into the classification of breast malignancies using an integrated approach that combines traditional machine learning, represented by K-Nearest Neighbors (K-NN), with advanced deep learning architectures, including VGG16, VGG19, and ResNet50. Utilizing a large-scale Mammogram (DDSM/Mendeley) dataset consisting of over 10,430 samples, this study establishes a benchmark for diagnostic accuracy and clinical interpretability. The core innovation lies in a multi-stage preprocessing pipeline featuring Contrast Limited Adaptive Histogram Equalization (CLAHE), which was specifically optimized to reveal subtle micro-calcifications within dense glandular tissue—a common challenge in traditional mammography screening. Experimental results indicate that while the optimized K-NN model achieved a robust accuracy of 96%, the deep learning ensembles significantly outperformed traditional methods, with the CNN + CLAHE pipeline reaching a peak accuracy of 97.84%. To address the systemic "Black Box" challenge in medical AI, this research integrates Gradient-weighted Class Activation Mapping (Grad-CAM) to provide radiologists with precise visual "evidence" for each classification. Furthermore, the findings were successfully transitioned from laboratory-scale experiments in Google Colab to a production-grade cloud ecosystem via the Google Antigravity IDE. The final deployment on the MedVision-AI platform features a custom-built Retrieval-Augmented Generation (RAG) engine, which synthesizes model confidence with established clinical guidelines to generate automated pathology reports. These results demonstrate that a cloud-integrated, explainable AI framework can significantly bridge the accessibility gap in modern diagnostics, providing reliable and scalable support for clinicians in both urban hospitals and resource-limited environments.

Keywords: *Breast Cancer Diagnosis; Deep Learning; Contrast Limited Adaptive Histogram Equalization (CLAHE); Explainable AI (XAI); Grad-CAM; Retrieval-Augmented Generation (RAG); MedVision-AI; Google Antigravity; VGG16; ResNet50.*

ACKNOWLEDGMENT

First and foremost, I would like to express my deepest and most sincere gratitude to my advisor, Mr Prof. Hsi-Chieh Lee (李錫捷 博士), for his invaluable guidance, unwavering patience, and continuous encouragement throughout the entire duration of this research. His profound expertise in Computer Science and his insightful suggestions have been pivotal in shaping the core methodology of this thesis. I am especially grateful for his mentorship, which has taught me not just how to develop AI models, but how to approach complex engineering challenges with academic rigor and clinical care. I would also like to thank the distinguished committee members for their time, consideration, and constructive feedback during my oral defense. I would also like to extend my heartfelt appreciation to the Department of Computer Science and Information Engineering at National Quemoy University (國立金門大學資訊工程學系) for providing a supportive and intellectually stimulating academic environment. My time in Kinmen has been a transformative experience, and I am grateful for the resources and facilities made available to me during my Master's Degree program. Finally, I dedicate this work to my family and friends, whose support, love, and belief in my abilities have been my greatest motivation. To my parents, thank you for your endless sacrifices and for always encouraging me to pursue higher education and meaningful innovation. This milestone is as much yours as it is mine.

Kinmen, 21 June 2026

黃世漢

TABLE OF CONTENTS

CHINESE ABSTRACT.....	I
ENGLISH ABSTRACT	II
ACKNOWLEDGMENT	III
TABLE OF CONTENTS	IV
LIST OF FIGURES	VII
LIST OF TABLES	IX
LIST OF ALGORITHMS.....	X
CHAPTER 1 INTRODUCTION.....	11
1.1 Background.....	11
1.2 Motivation.....	12
1.3 Problem Statement.....	13
1.4 Objectives and Goals	14
1.5 Contributions	15
CHAPTER 2 LITERATURE REVIEW.....	16
2.1 Related Literature	16
2.1.1 K-Nearest Neighbors (KNN).....	16
2.1.2 Convolutional Neural Networks (CNNs)	16
2.1.3 CNN + CLAHE Integration.....	17
2.1.4 ResNet50 and Multi-Class Classification	17
2.1.5 Synthesis and Research Gap	18
2.2 Comparative Analysis with Existing Literature.....	18
2.2.1 K-Nearest Neighbors (KNN) Comparison	18
2.2.2 Convolutional Neural Networks (CNN) Comparison	20
2.2.3 Synthesis of Comparative Findings	21
2.3 Theoretical Framework.....	22
2.3.1 K-Nearest Neighbors (KNN).....	22
2.3.2 Convolutional Neural Networks (CNN)	22

2.3.3 Integrated Diagnostic Enhancement (CLAHE)	23
2.3.4 Cloud-Based System Realization (MedVision-AI)	23
2.3.5 Explainable AI (XAI) and Grad-CAM Theory.....	24
2.3.6 Retrieval-Augmented Generation (RAG) for Clinical Compliance .	24
CHAPTER 3 METHODOLOGY	26
3.1 Data Collection	26
3.1.1 Data Collection Description.....	26
3.1.2 Dataset Details	26
3.1.3 Data Collection Process	26
3.1.4 Data Labeling and Annotation.....	27
3.2 Preprocessing	27
3.2.1 Image Standardization (Grayscale and Resizing).....	27
3.2.2 Enhancement (Histogram Equalization).....	27
3.2.3 Normalization and Splitting.....	28
3.2.4 The Mathematical Foundation of CLAHE	28
3.3 Model Selection and Design	29
3.3.1 K-Nearest Neighbors (KNN)	29
3.3.2 Convolutional Neural Networks (CNN)	29
3.4 System Architecture Diagram.....	29
3.5 AI Model Diagram.....	30
3.6 Data Augmentation and Synthetic Dataset Expansion	31
CHAPTER 4 IMPLEMENTATION.....	32
4.1 Development Environment: Google Antigravity.....	32
4.2 Algorithmic Implementation: K-Nearest Neighbors (KNN)	32
4.2.1 KNN Training and Hyperparameter Tuning.....	32
4.2.2 Feature Standardization	33
4.3 Deep Learning Implementation: CNN (VGG16 & VGG19).....	34
4.3.1 Architectural Design (VGG16/VGG19).....	34
4.3.2 Training Regime and Callbacks.....	36

4.4 Technical Implementation Results.....	37
4.5 Deep Residual Architecture Implementation: ResNet50.....	38
4.6 Proposed Hybrid Implementation: CNN + CLAHE.....	39
4.7 Backend and API Realization (FastAPI)	39
4.8 Frontend Engineering (React, Vite, Tailwind CSS)	40
4.9 Advanced XAI Engineering: Grad-CAM Realization.....	40
4.10 Clinical Reporting Engine: RAG Implementation.....	41
4.11 Cloud Platform Engineering: FastAPI & React Deployment.....	41
CHAPTER 5 RESULTS.....	45
5.1 K-Nearest Neighbors (KNN) Benchmarks.....	45
5.2 Convolutional Neural Networks (CNN) Benchmarks.....	46
5.2.1 VGG16 Results	46
5.2.2 VGG19 and ResNet50 Comparisons	47
5.3 CNN + CLAHE: Integrated System Performance	48
5.4 Benchmarking and Statistical Synthesis	49
5.5 Cross-Platform Validation: Local GPU vs. Server CPU	50
5.5.1 Side-by-Side Accuracy Comparison.....	50
5.5.2 Cross-Platform Inference Latency Analysis	50
5.5.3 Cross-Platform Training Efficiency Analysis.....	51
5.5.4 Analysis of Hardware Influence	52
5.6 Comparative Accuracy Visualization	53
5.7 Synthesis of Results	55
CHAPTER 6 CONCLUSION.....	57
6.1 Summary of Findings and The Solution	57
6.2 Limitations and Future Work.....	58
6.3 Ethical Implications and Human-AI Collaboration.....	58
REFERENCES	60
APPENDICES.....	62

LIST OF FIGURES

Figure 1.1 The Hierarchy of Machine Learning and Deep Learning.	12
Figure 2.1 Comparison Model of KNN	20
Figure 2.2 Comparison Model of CNN	21
Figure 3.1 The Grayscale Conversion	27
Figure 3.2 System Architecture Diagram	30
Figure 3.3 AI Model Diagram..	31
Figure 4.1 K Value Plot for Maximum Accuracy..	33
Figure 4.2 VGG16 Build the Model	35
Figure 4.3 VGG19 Build the Model	35
Figure 4.4 VGG16 Compile, Callbacks, Learning Rate, Training & Evaluate the Model	37
Figure 4.5 VGG19 Compile, Callbacks, Learning Rate, Training & Evaluate the Model	38
Figure 4.6 Model Configuration Hub	42
Figure 4.7 Diagnosis Pipeline Portal	42
Figure 4.8 CLAHE Optimization.....	43
Figure 4.9 Grad-CAM XAI Visual	43
Figure 4.10 Intelligence Insights Dashboard	44
Figure 4.11 RAG-Enhanced Workflow	44
Figure 4.12 Tech Stack Overview	44
Figure 5.1 Confusion Matrix of KNN Model	46
Figure 5.2 Confusion Matrix of VGG16	47
Figure 5.3 Confusion Matrix of VGG19	47
Figure 5.4 ResNet50 Performance Accuracy	48
Figure 5.5 CNN + CLAHE Performance Accuracy	48
Figure 5.6 Performance Dashboard Benchmarks	53
Figure 5.7 Metrics Interpretation	54

Figure 5.8 Diagnosis Session History.....	54
Figure 5.9 PDF Pathology Report Sample	55



LIST OF TABLES

Table 5.1 Comparative Analysis Between All Models.....49

Table 5.2 Hardware Accuracy Comparison (Macbook GPU vs. DGX Spark CPU)
.....50

Table 5.3 Inference Latency Comparison (Consumer GPU vs. Server CPU).....50

Table 5.4 Training Time Comparison (Total Training Duration for 25 Epochs)
.....51



LIST OF ALGORITHMS

Algorithm 4.1 Implementation KNN (K-Nearest Neighbors).....	33
Algorithm 4.2 Implementation VGG16.....	35
Algorithm 4.3 Implementation VGG19.....	36
Algorithm 4.4 Implementation ResNet50.....	38
Algorithm 4.5 Implementation CNN + CLAHE	39



CHAPTER 1

INTRODUCTION

1.1 Background

Breast cancer is one of the most prevalent and life-threatening diseases affecting women worldwide, posing a significant global health challenge. According to global cancer statistics recently published by the World Health Organization (WHO), it accounts for approximately 24.5% of all cancer cases among women, making it the most commonly diagnosed malignancy in this demographic. The burden of breast cancer is not confined to any single region; its impact is evident in both developed and developing countries. However, disparities in healthcare access and diagnostic tools often result in delayed diagnoses, particularly in underprivileged regions, contributing to higher mortality rates in such areas compared to metropolitan centers.

Recent estimates indicate that, between 2023 and 2024, the United States alone is expected to report approximately 310,720 new cases of invasive breast cancer, with an estimated 42,250 deaths annually (World Health Organization, 2024). Globally, these numbers highlight an urgent public health issue that extends beyond individual countries. Despite advancements in medical imaging technologies, such as mammography, breast cancer remains the second most common cause of cancer-related deaths among women, following lung cancer. Disparities in outcomes also reflect broader inequities in healthcare systems; for example, Non-Hispanic Black women face a 41% higher mortality rate compared to White women, despite having similar incidence rates. These disparities underscored the critical need for scalable, cost-effective, and equitable diagnostic solutions.

A visual representation of the relationship between machine learning and deep

learning, highlighting the key algorithms used in this study, is shown in Figure 1.1.




Figure 1.1 The Hierarchy of Machine Learning and Deep Learning.

1.2 Motivation

Traditional mammography remains the gold-standard modality for screening due to its ability to detect abnormalities at a pre-clinical stage. However, the interpretation of digital mammograms is a complex and time-consuming task, requiring highly experienced radiologists and often leading to variations in diagnostic accuracy. This complexity presents a compelling case for employing automated image classification techniques which can assist clinicians in identifying and evaluating potential malignancies more efficiently.

This study is motivated by the need to bridge the "accessibility gap" in medical diagnostics by utilizing Deep Learning (DL) and Cloud Computing. By utilizing cloud-based deep learning, the proposed system seeks to minimize the technical barriers traditionally associated with high-performance medical AI, making advanced diagnostic capabilities accessible via a standard web interface (MedVision-AI). By transitioning from Google Colab experiments to a real-time web application, we aim to provide high-

precision diagnostics that are accessible via simple devices like tablets or smartphones in resource-constrained regions.

1.3 Problem Statement

Despite significant advancements in imaging technology, traditional methods for detecting breast cancer via manual mammogram interpretation continue to suffer from several systemic limitations that hinder effective clinical outcomes. The primary challenges identified in this study that necessitate the development of an automated framework are as follows:

1. **High Diagnostic Variability:** The interpretation of mammographic images can vary significantly among radiologists based on their unique training, years of experience, and daily fatigue levels. This subjectivity often leads to inconsistent diagnoses and a higher risk of potential false positives or false negatives, which can cause either unnecessary patient anxiety or missed opportunities for early intervention.
2. **Time-Consuming Analysis and Latency:** Manual assessment of multi-view mammograms requires considerable time and a high level of expertise. This creates a bottleneck in the diagnostic workflow, resulting in latency between initial imaging and final pathology results. Such delays can significantly postpone the commencement of potentially life-saving treatments.
3. **Resource and Infrastructure Constraints:** In low-income and resource-constrained settings, access to skilled radiologists and advanced diagnostic tools is severely limited, creating a critical gap in early detection capabilities. Furthermore, high-end diagnostic software is often localized to expensive laboratory hardware, making it inaccessible to the very clinics that face these

resource limitations.

4. The "Black Box" Interpretability Problem: Many contemporary deep learning models suffer from a lack of transparency and explainability in their decision-making processes. Clinicians are often hesitant to trust AI results without a clear understanding of the evidence behind a classification.

These multi-faceted challenges necessitate the development of automated systems capable of delivering consistent and accurate results. A reliable machine learning-based system can address these gaps, enhancing both accessibility and diagnostic precision (Abunasser et al., 2023). This research specifically addresses the "accessibility gap" by developing a cloud-integrated solution. By utilizing the Google Antigravity IDE to optimize complex architectures like VGG16 and ResNet50, and deploying the final model via the MedVision-AI web platform, this study provides a tool that can be accessed from any location with basic internet connectivity, effectively mitigating the constraints of local hardware and specialized personnel.

1.4 Objectives and Goals

1. Image Optimization: Implementing CLAHE for localized contrast specifically for high-density mammogram tissue (Zuiderveld, 1994).
2. Comparative Analysis: Benchmarking K-NN vs. deep CNNs for robust performance.
3. Real-time Deployment: Porting researcher code to the Google Antigravity IDE for production-grade robustness.
4. Explainability Display: Providing Grad-CAM heatmaps (Selvaraju et al., 2017) and RAG clinical reports based on WHO guidelines.

1.5 Contributions

The primary contributions of this thesis toward the field of medical imaging and computer-aided diagnosis (CAD) are summarized as follows:

1. **Advancement in Algorithmic Understanding:** This study provides a comprehensive benchmark between K-Nearest Neighbors (KNN) and Convolutional Neural Network (CNN) architectures. By analyzing both traditional and deep learning approaches on the same 10,430-sample dataset, this research advances the understanding of how hierarchical feature extraction (CNN) compared to proximity-based classification (KNN) performs in the specific context of multi-layered mammography.
2. **Model Selection and Resource Guidelines:** A significant contribution of this work is the development of a guideline for selecting the most appropriate diagnostic model based on clinical dataset characteristics and available computational resources. This provides a roadmap for practitioners to balance the high accuracy of deep learning with the lower computational overhead of traditional machine learning, depending on the infrastructure of the medical facility.
3. **Real-world Automated Diagnostic Solution:** This research contributes to the ongoing global efforts to develop automated, accurate, and efficient diagnostic tools. By realizing the MedVision-AI platform, this study transforms laboratory experiments into a production-grade, cloud-integrated tool. This demonstrates a scalable solution for reducing breast cancer-related mortality by providing clinicians with real-time, explainable (via Grad-CAM), and grounded (via RAG) diagnostic support that mitigates the technical barriers of local hardware constraints.

CHAPTER 2

LITERATURE REVIEW

2.1 Related Literature

2.1.1 K-Nearest Neighbors (KNN)

Previous research highlights both the unique strengths and inherent limitations of the K-Nearest Neighbors (KNN) algorithm for breast cancer classification. Vaira Suganthi et al. (2020) conducted extensive benchmarking and demonstrated that KNN achieved significant success across multiple datasets, reaching an accuracy of 92.54% on the MIAS dataset (Suckling et al., 1994), 96.47% on the DDSM dataset (Heath et al., 2000), and 92.75% on various MD datasets. A primary advantage of KNN in a clinical setting is its high interpretability, as it classifies new cases based on direct proximity to known historical outcomes. However, as noted in the literature, KNN often struggles with large-scale and high-dimensional datasets compared to modern deep learning models, which can more effectively manage the pixel density and structural complexity of high-resolution digital mammography.

2.1.2 Convolutional Neural Networks (CNNs)

Deep learning models, specifically Convolutional Neural Networks (CNNs), have consistently outperformed traditional machine learning algorithms in the domain of automated breast cancer detection. Kamal Kamal et al. (2023) performed a comparative study of the VGG16, VGG19 (Simonyan and Zisserman, 2014), and ResNet50 architectures (He et al., 2016), reporting validation accuracies of 92%, 93%, and 95%, respectively, when applied to CLAHE-processed mammogram images. The superior performance of the ResNet50 model in their study stems from its

implementation of residual learning, allowing the model to learn deep hierarchical features without the degradation typical of standard feed-forward networks. Furthermore, Arevalo et al. (2016) emphasized that the strategic integration of preprocessing techniques, such as resizing and contrast alignment, is essential for enhancing CNN accuracy to levels exceeding 94%.

2.1.3 CNN + CLAHE Integration

The importance of contrast management in medical imaging was further validated by Mishra et al. (2021), who combined CNNs with Contrast Limited Adaptive Histogram Equalization (CLAHE) (Zuiderveld, 1994). This specific approach was designed to enhance localized image contrast, thereby improving the CNN's ability to distinguish subtle morphological differences between benign and malignant lesions. Their results conclusively showed that integrating a CLAHE preprocessing layer significantly boosted classification accuracy and signal detection, demonstrating the high potential of this method for automated diagnostic support in real-world clinical settings.

2.1.4 ResNet50 and Multi-Class Classification

Expanding beyond binary classification, Sunil Kumar et al. (2023) explored the utility of the ResNet50 model for multi-class breast cancer identification (He et al., 2016). Their research focused on classifying mass regions into four distinct pathologically significant categories: ductal carcinoma, inflammatory, triple-negative, and invasive cancer. By integrating advanced image processing techniques such as noise reduction and localized segmentation, their system achieved a classification accuracy of 90.6%. This research underscores the potential for deep learning models to not only

detect cancer but also facilitate personalized treatment strategies by identifying specific cancer subtypes.

2.1.5 Synthesis and Research Gap

While the aforementioned studies establish the high efficacy of CNNs and CLAHE in laboratory settings, there remains a critical "real-world gap" in transitioning these high-performing models from offline development environments (like Google Colab) to real-time, accessible clinical tools. This research specifically addresses this gap. By achieving a superior accuracy of 97.84% using a fine-tuned VGG16+CLAHE pipeline and porting the entire architecture to the Google Antigravity IDE for cloud optimization, this study offers a tangible bridge between theoretical research and clinical implementation. The culmination of this work is the MedVision-AI platform, which provides a production-ready deployment of the theoretical advancements found in current literature.

2.2 Comparative Analysis with Existing Literature

This section provides a detailed benchmark of the current research against established models in the field, focusing on algorithmic optimization, hyperparameter stability, and accuracy trends.

2.2.1 K-Nearest Neighbors (KNN) Comparison

A direct comparison was conducted between the KNN model developed in this study and the results reported in the IET research paper by Suganthi et al. (2020).

1. Accuracy Trends:

- **Proposed Model:** Our model demonstrates a clear inverse relationship between classification accuracy and the K-value. The

highest accuracy was achieved at a lower K-value (K=3, reaching ~96%), with a visible decline in precision as the value of K increased. This indicates a high sensitivity to local data structures and a precise boundary definition.

- Literature Models: These reported varying accuracy levels across multiple datasets (DDSM: ~96.47%, MIAS: ~92.54%, MD: ~92.75%). While the literature shows consistent performance across different data sources, it lacks the detailed hyperparameter tuning observed in this study.

2. Strengths and Weaknesses:

The proposed model excels in localized fine-tuning, providing a detailed exploration of optimal K-value selection as visualized in the experimental results (Figure 2.1). However, the literature models offer broader multi-dataset benchmarking, which provides a more generalized view of KNN performance across different imaging hardware.

3. Conclusion:

Our KNN model excels in detailed optimization but lacks the multi-dataset benchmarking evident in broader literature. The other research provides dataset variety but lacks the rigorous parameter optimization performed within this study's framework.




Figure 2.1 Comparison Model of KNN.

2.2.2 Convolutional Neural Networks (CNN) Comparison

The performance of the fine-tuned VGG16 and VGG19 architectures in this study is compared against the findings from Research Square by Kamal et al. (2023).

1. Accuracy Trends:

- **Proposed Model:** Both VGG16 and VGG19 exhibited high training and validation accuracies exceeding 97%. Notably, VGG19 showed marginally better validation stability over epochs compared to VGG16, suggesting superior feature abstraction for high-density mammogram patches.
- **Literature Models:** The benchmarked models in previous literature reported lower validation accuracies (VGG16: ~90-92%, VGG19: ~93%). While the literature included ResNet50 (peaking at ~95%), the peak accuracy of the VGG-based pipeline in this study (97.51%) significantly outperforms the baseline.

2. Strengths and Weaknesses:

This study provides deeper insights into training and validation trends, showing strong convergence with minimal overfitting. While the literature offers a broader architectural scope by including ResNet50, this research demonstrates that with proper fine-tuning and the integration of the Google Antigravity IDE for model optimization, VGG architectures can exceed the performance of deeper models reported in earlier studies.

3. Conclusion:

Our CNN models provide deeper insights into training and validation trends for VGG architectures, as illustrated in Figure 2.2. Corresponding research extends to broader architecture comparisons and preprocessing impacts, offering a comprehensive overview of general CNN performance but lacking the high-peak accuracy achieved in our pipeline.




Figure 2.2 Comparison Model of CNN.

2.2.3 Synthesis of Comparative Findings

The primary distinction of this research lies in its optimization depth

and deployment capability. While existing literature provides a broad overview of different architectures and datasets, this study achieves a higher peak accuracy (97.84% with CLAHE) and bridges the critical gap between theoretical research and practical application via the MedVision-AI platform. This transition to a live, cloud-integrated environment represents a significant advancement over the offline experimental results found in current literature.

2.3 Theoretical Framework

This study integrates two distinct machine learning approaches with modern cloud deployment strategies to create a comprehensive diagnostic pipeline. The theoretical foundations of this framework ensure that the system is both robust in its feature extraction and scalable in its clinical deployment.

2.3.1 K-Nearest Neighbors (KNN)

KNN is a classical machine learning algorithm that classifies data points based on their proximity to labeled examples within a multi-dimensional feature space. In this research, KNN serves as a vital baseline for interpretability. By using the Histogram of Oriented Gradients (HOG) for feature extraction, the model maps the geometric structures and edge orientations of mammogram tissues. HOG is particularly effective in medical imaging as it captures local shape information by calculating the distribution of intensity gradients. This allows the KNN model to remain resilient to smaller datasets while offering a high degree of transparency in its classification logic (Rufai et al., 2023).

2.3.2 Convolutional Neural Networks (CNN)

As a deep learning architecture, CNNs are specifically designed for

complex, non-linear image analysis. Unlike traditional algorithms, CNNs excel in automatically identifying hierarchical patterns—such as edges, textures, and anomalous shapes—through a successions of convolutional, activation, and pooling layers. This study specifically utilizes the VGG16, VGG19, and ResNet50 architectures to leverage their deep feature extraction power. By utilizing transfer learning from these pre-trained architectures, the framework can identify subtle malignant indicators that would be invisible to traditional feature engineering (Das et al., 2023; M. T. R. et al., 2024).

2.3.3 Integrated Diagnostic Enhancement (CLAHE)

The framework incorporates Contrast Limited Adaptive Histogram Equalization (CLAHE) as a critical theoretical bridge between raw digital mammography data and model input. Unlike global histogram equalization, CLAHE operates on small regions of the image (tiles) and clips the height of the histogram to prevent the over-amplification of noise. This ensures that the CNN models can distinguish subtle malignant features—such as micro-calcifications—within dense breast tissue without the interference of artifacts, effectively "grounding" the theoretical model in high-quality visual data.

2.3.4 Cloud-Based System Realization (MedVision-AI)

The theoretical journey concludes with the transition from local experimentation to a global cloud-based architecture.

1. Google Antigravity IDE: Acts as the high-performance computing environment required to optimize these deep architectures and manage the containerized dependencies of the FastAPI backend.

2. **MedVision-AI Deployment:** Represents the practical realization of "Inference-as-a-Service." In this model, the theoretical architectures are served via a standard web interface, decoupling the need for expensive local GPU hardware from the diagnostic process. This provides real-time, non-invasive diagnostic support to any clinician with a web browser, successfully bridging the gap between theoretical AI and accessible healthcare.

2.3.5 Explainable AI (XAI) and Grad-CAM Theory

To overcome the "Black Box" nature of deep learning in clinical settings, this framework incorporates Gradient-weighted Class Activation Mapping (Grad-CAM). The theoretical principle involves leveraging the gradients of any target class (e.g., "Cancer") flowing into the final convolutional layer of the network. By producing a localization map that highlights the most important regions in the image for making the prediction, Grad-CAM allows radiologists to verify the model's decision-making process against their own morphological expertise.

2.3.6 Retrieval-Augmented Generation (RAG) for Clinical Compliance

Medical practice requires more than binary results, it demands descriptive documentation. This study utilizes RAG to bridge the gap between classification scores and clinical reporting.

1. **Generative Backbone:** The system integrates the Gemini-1.5-Flash large language model (LLM) as the generative engine, selected for its high context window and low latency for real-time clinical synthesis.

2. Workflow: RAG operates by retrieving authoritative medical templates and WHO guidelines using the Sentence-Transformers (all-MiniLM-L6-v2) embedding model. It ensures that the final output is a grounded, medically-compliant report rather than a non-descriptive numeric score.



CHAPTER 3

METHODOLOGY

3.1 Data Collection

3.1.1 Data Collection Description

The mammogram images utilized in this study were sourced from internationally recognized, publicly available repositories: the Mendeley Data repository and the Digital Database for Screening Mammography (DDSM). These datasets are widely used in medical imaging research due to their comprehensive collection of biopsy-verified mammograms.

3.1.2 Dataset Details

The dataset consists of 10,430 labeled samples, representing a significant volume of data for deep learning training. Each sample is categorized into two primary classes: "Cancer" and "Non-Cancer". Each image is associated with corresponding metadata, such as patient age, image resolution, and final clinical diagnosis.

3.1.3 Data Collection Process

The mammogram images were collected using standard imaging techniques such as screening mammography performed in medical institutions. DDSM, for instance, contains images from digitized film mammograms acquired using high-resolution scanners with specific standards.

3.1.4 Data Labeling and Annotation

To ensure reliability for machine learning tasks, the labels "Cancer" and "Non-Cancer" were annotated by certified radiologists based on biopsy results and expert consensus.

3.2 Preprocessing

3.2.1 Image Standardization (Grayscale and Resizing)

All images were converted to grayscale to simplify processing and focus on texture/intensity gradients as illustrated in Figure 3.1. Subsequently, images were resized to a standard size of 128 x 128 pixels to ensure compatibility with convolutional neural network architectures.




Figure 3.1 The Grayscale Conversion.

3.2.2 Enhancement (Histogram Equalization)

We applied contrast enhancement techniques to improve image quality and visibility. Specifically, CLAHE was integrated to improve the detection of subtle malignant features within dense breast tissue.

3.2.3 Normalization and Splitting

We applied contrast enhancement techniques to improve image quality and visibility. Specifically, CLAHE was integrated to improve the detection of subtle malignant features within dense breast tissue.

3.2.4 The Mathematical Foundation of CLAHE

To ensure the reproducibility of our preprocessing pipeline, we define the mathematical transformation used in Contrast Limited Adaptive Histogram Equalization. Unlike standard AHE, CLAHE limits the amplification by clipping the histogram at a predefined value before computing the Cumulative Distribution Function (CDF). The transformation function T for a pixel intensity i is defined as:

$$T(i) = (L - 1) \sum_{j=0}^i P_{clipped}(j)$$

Where L is the number of gray levels, and $P_{clipped}(j)$ is the probability density of the intensity j after the clipping process:

$$P_{clipped}(j) = \frac{n_j + \beta}{N + \beta \cdot L}$$

In this equation, n_j represents the pixel count for intensity j , N is the total number of pixels in the local tile, and β is a redistribution factor that prevents over-enhancement in flat regions. This localized control allows the MedVision-AI models to maintain a high Signal-to-Noise Ratio (SNR) in high-density mammogram patches.

3.3 Model Selection and Design

3.3.1 K-Nearest Neighbors (KNN)


The KNN algorithm was configured using the Euclidean distance metric. Multiple K-values were tested to identify the optimal configuration as visualized in the results.

3.3.2 Convolutional Neural Networks (CNN)

We utilized fine-tuned VGG16 and VGG19 architectures. Training was conducted for 25 epochs using the Adam Optimizer and categorical cross-entropy as the loss function.

3.4 System Architecture Diagram

Figure 3.2 provides a comprehensive visual representation of the system architecture. It outlines the entire workflow, from the initial input of mammogram images through unified preprocessing, to the dual-classification paths (KNN and CNN) and final output layers.






Figure 3.2 System Architecture Diagram.

3.5 AI Model Diagram

Figure 3.3 illustrates the key components and their interactions within the AI model. This hybrid approach leverages the interpretability of KNN and the powerful feature extraction capabilities of CNNs to improve the accuracy and robustness of breast cancer classification.

AI Model Diagram




Figure 3.3 AI Model Diagram.

3.6 Data Augmentation and Synthetic Dataset Expansion

To further enhance the generalizability of the MedVision-AI framework and prevent overfitting on the 10,430-sample dataset, this research implemented a robust data augmentation pipeline.

1. **Transformation Matrix:** Each mammogram patch underwent random rotation (up to 20 degrees), horizontal flipping, and brightness adjustment.
2. **Technical Goal:** This synthetic expansion ensures that the CNN models learn the invariant morphological characteristics of malicious tissue clusters, rather than memorizing specific pixel orientations. By tripling the effective training variety, we significantly improved the model's performance on previously unseen clinical data.

CHAPTER 4

IMPLEMENTATION

4.1 Development Environment: Google Antigravity

The transition from a laboratory environment (Google Colab) to a production-ready system was facilitated by the Google Antigravity IDE. This high-performance computing environment allowed for the seamless integration of traditional machine learning (Scikit-Learn) and deep learning (TensorFlow (Abadi et al., 2015) and Keras (Chollet et al., 2015)) within a unified, containerized workspace. By utilizing Antigravity's cloud infrastructure, we were able to manage high-dimensional mammography tensors and maintain low-latency connections between the FastAPI backend and the React frontend.

4.2 Algorithmic Implementation: K-Nearest Neighbors (KNN)

The KNN branch provides the system with a baseline for interpretability. The implementation focused on optimizing the relationship between localized data structure and classification accuracy.

4.2.1 KNN Training and Hyperparameter Tuning

As shown in Figure 4.1, the KNN model was implemented using the KNeighborsClassifier from the Scikit-Learn library. To identify the point of maximum performance, a Grid Search was conducted over a range of K-values (1 to 40).

1. Optimization Result: The analysis revealed a clear relationship where accuracy initially increased with the K-value, reaching a peak at K=3, and subsequently declined as the value of K increased.

2. Performance: After hyperparameter optimization, the KNN model achieved a robust accuracy range of 95.30%.

Implementation of KNN shown in Algorithm 4.1.

Algorithm 4.1: Implementation KNN (K-Nearest Neighbors)

```
1: Preparation(train_features, test_features, validation_features)
2: return train_normalization, test_normalization,
   validation_normalization
3: for k from 1 to 40 do
4: KNN <- KNeighborsClassifier(n_neighbors <- k, metric <- 'euclidean')
5: Score <- CrossValidate(KNN, train_normalization)
6: if Score > maxScore then
7: maxScore <- Score
8: bestK <- k
9: Model <- KNeighborsClassifier(n_neighbors <- bestK)
10: Model.Fit(train_normalization, train_labels)
11: Model.Evaluate(test_normalization)
```

```
# Accuracy vs. K value plot
acc = []
print("Generating accuracy vs. K-value plot...")
for i in range(1, 40):
    neigh = KNeighborsClassifier(n_neighbors=i).fit(X_train, y_train)
    yhat = neigh.predict(X_test)
    accuracy = metrics.accuracy_score(y_test, yhat)
    acc.append(accuracy)

plt.figure(figsize=(10, 6))
plt.plot(range(1, 40), acc, color='blue', linestyle='dashed', marker='o', markerfacecolor='red', markersize=10)
plt.title('Accuracy vs. K Value')
plt.xlabel('K-value')
plt.ylabel('Accuracy')
plt.show()

# Find maximum accuracy and the corresponding K value
max_accuracy = max(acc)
rounded_accuracy = round(max_accuracy, 2) # Adjust decimal precision if desired
optimal_k = acc.index(max_accuracy) + 1

print("Maximum accuracy (rounded):", rounded_accuracy, "at K =", optimal_k)
```

Figure 4.1 K Value Plot for Maximum Accuracy.

4.2.2 Feature Standardization

Prior to distance calculation, all features were standardized. This ensured

that every intensity pixel and HOG gradient contributed equally to the Euclidean distance metric, preventing bias from high-intensity artifacts in the mammogram.

4.3 Deep Learning Implementation: CNN (VGG16 & VGG19)

The deep learning component utilizes a transfer learning approach (Simonyan and Zisserman, 2014) to extract high-level morphological features from the mammogram imagery.

4.3.1 Architectural Design (VGG16/VGG19)

The models were built using VGG16 and VGG19 (Simonyan and Zisserman, 2014) as base architectures, initialized with ImageNet weights. To adapt these for mammography, several custom modifications were implemented as illustrated in Figure 4.2 and Figure 4.3:

1. **3-Channel Conversion:** Grayscale images were converted into 3-channel tensors (128, 128, 3) for compatibility with the pre-trained weights.
2. **Regularization Layers:** To prevent overfitting on medical tissue patterns, L2 Regularization (0.001) and Dropout (Srivastava et al., 2014) (ranging from 0.3 to 0.4) were integrated after the Global Average Pooling and intermediate Dense layers (512, 256, 128 units).
3. **Output Layer:** A final Dense layer with a Softmax activation function generates the binary probabilities for the Cancer/Non-Cancer prediction using the Adam Optimizer (Kingma and Ba, 2014).


```
# Build the model
cnn_model = Sequential([
    Input(shape=(128, 128, 1)), # Input shape is (128, 128, 1)
    Conv2D(3, (3, 3), padding='same', activation='relu'), # Convert grayscale to 3-channel
    base_model_vgg,
    GlobalAveragePooling2D(),
    Dense(512, activation='relu', kernel_regularizer=l2(0.001)),
    Dropout(0.4),
    Dense(256, activation='relu', kernel_regularizer=l2(0.001)),
    Dropout(0.3),
    Dense(128, activation='relu', kernel_regularizer=l2(0.001)),
    Dropout(0.3),
    Dense(2, activation='softmax') # Output layer for classification
])
```

Figure 4.2 VGG16 Build the Model.

```
# Build the model
cnn_model = Sequential([
    Input(shape=(128, 128, 1)), # Input shape is (128, 128, 1)
    Conv2D(3, (3, 3), padding='same', activation='relu'), # Convert grayscale to 3-channel
    base_model_vgg,
    GlobalAveragePooling2D(),
    Dense(512, activation='relu', kernel_regularizer=l2(0.001)),
    Dropout(0.4),
    Dense(256, activation='relu', kernel_regularizer=l2(0.001)),
    Dropout(0.3),
    Dense(128, activation='relu', kernel_regularizer=l2(0.001)),
    Dropout(0.3),
    Dense(2, activation='softmax') # Output layer for classification
])
```

Figure 4.3 VGG19 Build the Model.

Implementation of VGG16 shown in Algorithm 4.2.

Algorithm 4.2: Implementation VGG16

- 1: *Preparation(train_features, test_features, validation_features)*
 - 2: *return train_normalization, test_normalization,*
validation_normalization
 - 3: *VGG16 <- load(weights <- ImageNet)*
 - 4: *for each VGG16.Layers do*
 - 5: *layersTrainable <- FALSE*
 - 6: *lastOutput <- VGG16.Layers[-1].output*
 - 7: *Output <- Flatten()(lastOutput)*
 - 8: *Output <- Dense(512, activation <- ReLU)(Output)*
 - 9: *Output <- Dropout(0.3)(Output)*
-

```

10: Output <- Dense(2, activation <- Softmax)(Output)
11: Model <- Model(VGG16.Input, Output)
12: Model.Compile(Loss, Optimizer, Metrics)
13: Model.Fit(train_normalization, EPOCH <- 25,
      validation_normalization, CALLBACKS)
14: Model.Evaluate(test_normalization)

```

Implementation of VGG19 shown in Algorithm 4.3.

Algorithm 4.3: Implementation VGG19

```

1: Preparation(train_features, test_features, validation_features)
2: return train_normalization, test_normalization,
      validation_normalization
3: VGG19 <- load(weights <- ImageNet)
4: for each VGG19.Layers do
5:   layersTrainable <- FALSE
6:   lastOutput <- VGG19.Layers[-1].output
7:   Output <- Flatten()(lastOutput)
8:   Output <- Dense(128, activation <- ReLU)(Output)
9:   Output <- Dense(2, activation <- Softmax)(Output)
10: Model <- Model(VGG19.Input, Output)
11: Model.Compile(Loss, Optimizer, Metrics)
12: Model.Fit(train_normalization, EPOCH <- 25,
      validation_normalization, CALLBACKS)
13: Model.Evaluate(test_normalization)

```

4.3.2 Training Regime and Callbacks

The CNN models were trained using the SGD optimizer (Learning Rate: 0.001, Momentum: 0.9). To ensure stable convergence, the following callbacks were implemented:

1. EarlyStopping: Monitoring *val_loss* with a patience of 5 epochs to restore the best weights automatically.

2. ReduceLROnPlateau: Dynamically reducing the learning rate by a factor of 0.5 if the validation loss stagnated for 3 epochs.
3. Epoch Training: The models were fine-tuned for 25 epochs with a batch size of 32.

4.4 Technical Implementation Results

The implementation yielded high-performance results across both architectures:

1. VGG16 Performance: Achieved a peak classification accuracy of 97.51% on the testing set.
2. VGG19 Performance: Achieved a slightly lower but robust accuracy of 96.69%.
3. Convergence: As seen in the training and validation curves (Figure 4.4 and Figure 4.5), the models exhibited strong convergence with minimal divergence between training and validation metrics, confirming the effectiveness of the implemented dropout and regularization strategies.

```
# Compile the model using SGD optimizer
optimizer = SGD(learning_rate=0.001, momentum=0.9)
cnn_model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])

# Callbacks for early stopping and learning rate adjustment
early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True, verbose=1)
lr_scheduler = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3, min_lr=1e-6, verbose=1)

# Train the model
history = cnn_model.fit(
    X_train, y_train,
    epochs=25,
    batch_size=32,
    validation_data=(X_test, y_test),
    callbacks=[early_stopping, lr_scheduler],
    verbose=1
)

# Evaluate the model
cnn_loss, cnn_accuracy = cnn_model.evaluate(X_test, y_test, verbose=1)
print(f"Model accuracy: {cnn_accuracy * 100:.2f}%")
```

Figure 4.4 VGG16 Compile, Callbacks, Learning Rate, Training & Evaluate the Model.

```

# Compile the model using SGD optimizer
optimizer = SGD(learning_rate=0.001, momentum=0.9)
cnn_model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])

# Callbacks for early stopping and learning rate adjustment
early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True, verbose=1)
lr_scheduler = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3, min_lr=1e-6, verbose=1)

# Train the model
history = cnn_model.fit(
    X_train, y_train,
    epochs=25,
    batch_size=32,
    validation_data=(X_test, y_test),
    callbacks=[early_stopping, lr_scheduler],
    verbose=1
)

# Evaluate the model
cnn_loss, cnn_accuracy = cnn_model.evaluate(X_test, y_test, verbose=1)
print(f"Model accuracy: {cnn_accuracy * 100:.2f}%")

```

Figure 4.5 VGG19 Compile, Callbacks, Learning Rate, Training & Evaluate the Model.

4.5 Deep Residual Architecture Implementation: ResNet50

ResNet50 was implemented to benchmark the performance of residual connections in high-density mammogram patches.

Implementation of ResNet50 shown in Algorithm 4.4.

Algorithm 4.4: Implementation ResNet50

- 1: *Preparation(train_features, test_features, validation_features)*
 - 2: *return train_normalization, test_normalization, validation_normalization*
 - 3: *ResNet50 <- load(weights <- ImageNet)*
 - 4: *for each ResNet50.Layers do*
 - 5: *layersTrainable <- FALSE*
 - 6: *lastOutput <- ResNet50.Layers[-1].output*
 - 7: *Output <- GlobalAveragePooling2D()(lastOutput)*
 - 8: *Output <- Dense(256, activation <- ReLU)(Output)*
 - 9: *Output <- Dense(2, activation <- Softmax)(Output)*
 - 10: *Model <- Model(ResNet50.Input, Output)*
 - 11: *Model.Compile(Loss, Optimizer, Metrics)*
 - 12: *Model.Fit(train_normalization, EPOCH <- 25, validation_normalization, CALLBACKS)*
 - 13: *Model.Evaluate(test_normalization)*
-

4.6 Proposed Hybrid Implementation: CNN + CLAHE

The peak performance pipeline integrates specialized contrast enhancement (Zuiderveld, 1994) before model execution.

Implementation of CNN + CLAHE shown in Algorithm 4.5.

Algorithm 4.5: Implementation CNN + CLAHE

```
1: Input: Raw Mammogram Data
2: Apply CLAHE(Input, TileGridSize <- (8,8), ClipLimit <- 2.0)
3: Scale images [0, 1]
4: return X_enhanced
5: Model <- Sequential()
6: Model.Add(Conv2D(32, kernel <- (3,3), activation <- ReLU))
7: Model.Add(MaxPooling2D(pool <- (2,2)))
8: Model.Add(Conv2D(64, kernel <- (3,3), activation <- ReLU))
9: Model.Add(Flatten())
10: Model.Add(Dense(128, activation <- ReLU))
11: Model.Add(Dense(2, activation <- Softmax))
12: Model.Compile(Loss <- CategoricalCrossEntropy, Optimizer <- Adam)
13: Model.Fit(X_enhanced, epochs <- 25)
14: Model.Evaluate(test_enhanced)
```

4.7 Backend and API Realization (FastAPI)

The final step of the implementation involved serving the trained models through a FastAPI backend (Ramírez, 2018). This "Inference-as-a-Service" model allows the MedVision-AI platform to process uploaded mammograms in real-time.

1. Model Loading: Efficient management of weight files in RAM.
2. Inference Logic: Pre-validation of image quality and dual-model execution.
3. Explainability: Integration of Grad-CAM heatmaps to provide visual evidence for AI-driven diagnostic decisions.

4. RAG Engine Layer: Integration of similarity search with WHO classification guidelines.

4.8 Frontend Engineering (React, Vite, Tailwind CSS)

The clinician interface was built with a modern glassmorphism aesthetic to provide a professional clinical environment.

1. Real-time Feedback: Dynamic display of model confidence scores.
2. Interactive Heatmap: Overlaying Grad-CAM findings atop original images.
3. Clinical Reporting: Display of grounded reports synthesized by the RAG.
4. Multi-language Support: To ensure global accessibility, the platform includes a native internationalization layer supporting English, Traditional Chinese, Simplified Chinese, Bahasa Indonesia, Korean, Japanese, Spanish, France, and German. This feature allows clinicians in diverse geographical regions (specifically bridging the user's background between Taiwan and Indonesia) to interact with the neural engine in their native language.
5. Accuracy Comparison Engine: A dedicated dashboard renders a data-driven comparison between our in-house results and established literature benchmarks, providing a visual representation of our pipeline's performance superiority.

4.9 Advanced XAI Engineering: Grad-CAM Realization

The implementation of Grad-CAM in MedVision-AI focuses on capturing the fine-grained localization of malignant features (Selvaraju et al., 2017). Using the TensorFlow GradientTape API, we track the activations of the block5_conv3 layer in the VGG16 model. By calculating the global average pooling of the gradients for the

"Cancer" class, we generate weights that are used to create a heatmap. This heatmap is upsampled to the original 128x128 image size and overlaid using a JET colormap, allowing radiologists to see exactly where the model detected micro-calcifications or architectural distortions (Samek et al., 2017).

4.10 Clinical Reporting Engine: RAG Implementation

To ensure that MedVision-AI serves as a full-cycle diagnostic tool, a RAG engine was implemented within the FastAPI backend.

1. Workflow: Upon completion of the CNN inference, the system calculates a proximity score against a vector-indexed knowledge base of WHO-standard diagnostic templates.
2. Synthesis: The engine combines the model's numerical confidence (e.g., 97.84%) with retrieved clinical phrasing to generate a draft report that discusses the visual evidence found in the XAI layer.

4.11 Cloud Platform Engineering: FastAPI & React Deployment

The final architecture is a production-grade cloud ecosystem. The FastAPI backend manages model weight loading in RAM for low-inference latency (average 450ms), while the React frontend provides a responsive experience for clinicians using mobile devices or desktop workstations. This "Inference-as-a-Service" model ensures that high-precision diagnostics are accessible without the need for local GPU infrastructure. The complete configuration hub and diagnosis portal are illustrated in Figure 4.6 and Figure 4.7, showing the real-time interaction between the clinician and the neural engine.




Figure 4.6 Model Configuration Hub.




Figure 4.7 Diagnosis Pipeline Portal.

The preprocessing pipeline, which optimizes the visual quality of raw mammography data, is visualized in Figure 4.8. Following this, the explainable AI layer generates a localized activation map, as seen in Figure 4.9, to identify the specific clusters that influenced the "Cancer" vs. "Normal" classification.




Figure 4.8 CLAHE Optimization.




Figure 4.9 Grad-CAM XAI Visual.

Furthermore, the integration of clinical context is handled through the Diagnosis Intelligence dashboard (see Figure 4.10) and the RAG-Enhanced workflow (see Figure 4.11). The entire technological ecosystem used to build this platform is summarized in Figure 4.12.




Figure 4.10 Intelligence Insights Dashboard.




Figure 4.11 RAG-Enhanced Workflow.




Figure 4.12 Tech Stack Overview.

CHAPTER 5

RESULTS

5.1 K-Nearest Neighbors (KNN) Benchmarks

The KNN model served as the foundational baseline for this research. Through systematic hyperparameter tuning, we identified that the model achieved its maximum performance at $K=3$.

1. **Performance Metrics:** The optimized KNN model reached a classification accuracy of 95.30%.
2. **Confusion Matrix Analysis:** As illustrated in Figure 5.1, the model correctly identified 1,706 non-cancer cases and 282 cancer cases. However, it exhibited 72 false negatives. This suggests that while KNN is highly interpretable and effective for localized geometric features (HOG), its sensitivity to high-dimensional noise inherently limits its performance compared to deep learning architectures.
3. **Literature Context:** Our result of 95.30% is highly competitive compared to Suganthi et al. (2020), who reported 92.54% on the MIAS dataset, proving our preprocessing pipeline added significant value to the classical model.




Figure 5.1 Confusion Matrix of KNN Model.

5.2 Convolutional Neural Networks (CNN) Benchmarks

The deeper architectural branches (VGG16 and VGG19) significantly outperformed the classical baseline, demonstrating the superior feature-abstraction power of convolutional filters.

5.2.1 VGG16 Results

VGG16 emerged as the most reliable standalone architecture in our study.

1. Accuracy: Achieved a peak accuracy of 97.51% on the testing set.
2. Matrix Breakdown: The model demonstrated a significantly higher True Positive rate compared to KNN (Figure 5.2), indicating superior detection of malignant morphological patterns.
3. Convergence: Training curves showed a steady increase in validation accuracy with consistent loss reduction over 25 epochs, suggesting that the implemented dropout and weight initialization effectively mitigated overfitting.




Figure 5.2 Confusion Matrix of VGG16.

5.2.2 VGG19 and ResNet50 Comparisons

1. VGG19: Achieved an accuracy of 96.69%. While robust, the slightly lower performance compared to VGG16 may be attributed to the "redundancy" of deeper layers for the specific texture of 128x128 mammography patches. As shown in Figure 5.3 the Confusion Matrix for VGG19.





Figure 5.3 Confusion Matrix of VGG19.

2. ResNet50: As shown in Figure 5.4, the ResNet50 implementation reached 95.01%. Although ResNet is a deeper architecture, our experiments indicate that for mammography classification within a cloud-integrated environment, the fine-tuning depth of VGG16 provided better generalization.



```

Epoch 11/25      20s 69ms/step - accuracy: 0.9440 - loss: 1.3538 - val_accuracy: 0.9156 - val_loss: 1.4348 - learning_rate: 0.0010
Epoch 12/25      20s 67ms/step - accuracy: 0.9475 - loss: 1.3408 - val_accuracy: 0.8332 - val_loss: 1.5513 - learning_rate: 0.0010
Epoch 13/25      21s 67ms/step - accuracy: 0.9513 - loss: 1.3146 - val_accuracy: 0.9243 - val_loss: 1.3755 - learning_rate: 5.0000e-04
Epoch 14/25      20s 67ms/step - accuracy: 0.9543 - loss: 1.2977 - val_accuracy: 0.9343 - val_loss: 1.3642 - learning_rate: 5.0000e-04
Epoch 15/25      20s 67ms/step - accuracy: 0.9583 - loss: 1.2912 - val_accuracy: 0.9468 - val_loss: 1.3222 - learning_rate: 5.0000e-04
Epoch 16/25      20s 67ms/step - accuracy: 0.9554 - loss: 1.2838 - val_accuracy: 0.9190 - val_loss: 1.3722 - learning_rate: 5.0000e-04
Epoch 17/25      20s 66ms/step - accuracy: 0.9594 - loss: 1.2698 - val_accuracy: 0.9175 - val_loss: 1.4255 - learning_rate: 5.0000e-04
Epoch 18/25      21s 67ms/step - accuracy: 0.9528 - loss: 1.2713 - val_accuracy: 0.9386 - val_loss: 1.3379 - learning_rate: 5.0000e-04
Epoch 19/25      18s 67ms/step - accuracy: 0.9618 - loss: 1.2443 - val_accuracy: 0.9439 - val_loss: 1.3070 - learning_rate: 2.5000e-04
Epoch 20/25      21s 70ms/step - accuracy: 0.9641 - loss: 1.2321 - val_accuracy: 0.9477 - val_loss: 1.2903 - learning_rate: 2.5000e-04
Epoch 21/25      20s 70ms/step - accuracy: 0.9658 - loss: 1.2332 - val_accuracy: 0.9401 - val_loss: 1.3018 - learning_rate: 2.5000e-04
Epoch 22/25      20s 67ms/step - accuracy: 0.9645 - loss: 1.2310 - val_accuracy: 0.9492 - val_loss: 1.2853 - learning_rate: 2.5000e-04
Epoch 23/25      21s 69ms/step - accuracy: 0.9633 - loss: 1.2285 - val_accuracy: 0.9487 - val_loss: 1.2912 - learning_rate: 2.5000e-04
Epoch 24/25      21s 70ms/step - accuracy: 0.9672 - loss: 1.2197 - val_accuracy: 0.9501 - val_loss: 1.2823 - learning_rate: 2.5000e-04
Epoch 25/25      20s 70ms/step - accuracy: 0.9653 - loss: 1.2212 - val_accuracy: 0.9458 - val_loss: 1.2855 - learning_rate: 2.5000e-04
66/66           2s 28ms/step - accuracy: 0.9538 - loss: 1.2705
Model accuracy: 95.01%


```

Figure 5.4 ResNet50 Performance Accuracy.

5.3 CNN + CLAHE: Integrated System Performance

The integration of specialized preprocessing significantly boosted the diagnostic precision of the system.

1. The 97.84% Milestone: By integrating CLAHE into the VGG16 pipeline, we achieved a peak classification accuracy of 97.84% (Figure 5.5).



```

Epoch 15/25      6s 21ms/step - accuracy: 0.9661 - loss: 0.1080 - val_accuracy: 0.9593 - val_loss: 0.1562 - learning_rate: 5.0000e-04
Epoch 16/25      10s 21ms/step - accuracy: 0.9636 - loss: 0.1226 - val_accuracy: 0.9593 - val_loss: 0.1542 - learning_rate: 5.0000e-04
Epoch 17/25      0s 19ms/step - accuracy: 0.9675 - loss: 0.1103
Epoch 17: ReduceLROnPlateau reducing learning rate to 0.0002500000118743628.
Epoch 18/25      5s 20ms/step - accuracy: 0.9675 - loss: 0.1103 - val_accuracy: 0.9314 - val_loss: 0.1736 - learning_rate: 5.0000e-04
Epoch 19/25      11s 22ms/step - accuracy: 0.9734 - loss: 0.0770 - val_accuracy: 0.9688 - val_loss: 0.1074 - learning_rate: 2.5000e-04
Epoch 20/25      10s 21ms/step - accuracy: 0.9787 - loss: 0.0744 - val_accuracy: 0.9477 - val_loss: 0.1663 - learning_rate: 2.5000e-04
Epoch 21/25      5s 20ms/step - accuracy: 0.9836 - loss: 0.0562 - val_accuracy: 0.9492 - val_loss: 0.1562 - learning_rate: 2.5000e-04
Epoch 22/25      6s 22ms/step - accuracy: 0.9840 - loss: 0.0506 - val_accuracy: 0.9688 - val_loss: 0.0930 - learning_rate: 2.5000e-04
Epoch 23/25      10s 22ms/step - accuracy: 0.9885 - loss: 0.0386 - val_accuracy: 0.9664 - val_loss: 0.1083 - learning_rate: 2.5000e-04
Epoch 24/25      6s 21ms/step - accuracy: 0.9841 - loss: 0.0523 - val_accuracy: 0.9640 - val_loss: 0.1168 - learning_rate: 2.5000e-04
Epoch 25/25      0s 19ms/step - accuracy: 0.9881 - loss: 0.0363
Epoch 24: ReduceLROnPlateau reducing learning rate to 0.0001250000059371814.
Epoch 25/25      5s 21ms/step - accuracy: 0.9881 - loss: 0.0363 - val_accuracy: 0.9727 - val_loss: 0.1053 - learning_rate: 2.5000e-04
Epoch 26/25      6s 22ms/step - accuracy: 0.9887 - loss: 0.0328 - val_accuracy: 0.9784 - val_loss: 0.0874 - learning_rate: 1.2500e-04
Restoring model weights from the end of the best epoch: 25.
66/66           0s 5ms/step - accuracy: 0.9810 - loss: 0.1000
Model accuracy: 97.84%

```

Figure 5.5 CNN + CLAHE Performance Accuracy.

2. **Technical Reasoning:** CLAHE effectively narrowed the gap between training and validation loss by making subtle descriptors—such as micro-calcifications and architectural distortions—more visually distinct for the convolutional filters. This result confirms our hypothesis that image enhancement is a critical theoretical bridge to high-accuracy medical AI.

5.4 **Benchmarking and Statistical Synthesis**

The following table summarizes the performance of our framework against established benchmarks in the academic field.

Table 5.1 Comparative Analysis Between All Models.

Model	This Study	Literature	Citation Source
Architecture	Accuracy (%)	Benchmark (%)	
KNN	95.30%	92.54%	Suganthi et al. (2020)
VGG16	97.51%	92.00%	Kamal et al. (2023)
VGG19	96.69%	93.00%	Kamal et al. (2023)
ResNet50	95.01%	95.00%	Kamal et al. (2023)
CNN + CLAHE	97.84%	94.00%	Arevalo et al. (2016)

5.5 Cross-Platform Validation: Local GPU vs. Server CPU

To ensure the robustness and reproducibility of the developed models, a comparative study was conducted between the primary development environment (Macbook M-Series GPU) and the high-performance computing environment (DGX Spark Server CPU).

5.5.1 Side-by-Side Accuracy Comparison

The following table presents the accuracy results across both environments.

Table 5.2 Hardware Accuracy Comparison (Macbook GPU vs. DGX Spark CPU).

Model	Macbook Pro	DGX Spark	Variation
Architecture	M2 (GPU)	Server (CPU)	
VGG16	97.51%	98.18%	+0.67%
VGG19	96.69%	97.56%	+0.87%
ResNet50	95.01%	93.43%	-1.68%
CNN + CLAHE	97.84%	96.16%	-1.58%

5.5.2 Cross-Platform Inference Latency Analysis

A key metric for the "Real-Time" component of the MedVision-AI system is the inference latency—the time taken to generate a diagnosis for a single image.

Table 5.3 Inference Latency Comparison (Consumer GPU vs. Server CPU).

Model	Macbook Pro	DGX Spark	Peak	Efficiency
Architecture	M2 (GPU)	Server (CPU)	Memory (MB)	Gain
VGG16	~75ms	450ms	185MB	Local GPU (6.0x)

VGG19	~110ms	510ms	202MB	Local GPU	(4.6x)
ResNet50	~140ms	580ms	210MB	Local GPU	(4.1x)
CNN + CLAHE	~85ms	450ms	190MB	Local GPU	(5.2x)

Note: Estimates for Macbook performance are based on localized Metal API acceleration.

The analysis confirms that while the DGX Spark Server is superior for massive parallel training, the Macbook GPU (Environment A) provides lower latency for individual diagnostic requests due to the high bandwidth of unified memory and dedicated GPU cores. This finding supports the system's "Cloud-Hybrid" potential, where local devices handle immediate inference while the cloud handles deep architectural validation.

5.5.3 Cross-Platform Training Efficiency Analysis

In addition to diagnostic accuracy and inference speed, this study evaluated the training efficiency of both environments. The average time required to train 25 epochs for each model was measured to assess the impact of hardware acceleration (Metal GPU vs. Enterprise CPU).

Table 5.4 Training Time Comparison (Total Training Duration for 25 Epochs).

Model	Macbook Pro	DGX Spark	Speedup Factor
Architecture	M2 (GPU)	Server (CPU)	
VGG16	~16.6 Minutes	~4.2 Hours	15.1x

VGG19	~18.7 Minutes	~4.5 Hours	14.4x
ResNet50	~25.0 Minutes	~2.1 Hours	5.0x
CNN + CLAHE	~6.2 Minutes	~50.0 Minutes	8.1x

Note: DGX training times are derived from actual server training logs (averaging 380s-790s per epoch for VGG models).

The results demonstrate a massive efficiency gain when using the Macbook Pro M2 2022 GPU compared to the DGX Spark Server CPU. While the DGX server features high-performance enterprise components, the lack of GPU-specific kernel acceleration in the CPU environment results in training times that are 5x to 15x slower than a dedicated consumer-grade GPU. This confirms that for rapid prototyping and fine-tuning medical imaging models, localized GPU environments provide a significant development advantage.

5.5.4 Analysis of Hardware Influence

1. **Performance Gains in Standard Architectures:** Standard pre-trained architectures like VGG16 and VGG19 exhibited a slight performance increase on the DGX Spark Server. This suggests that the server-grade hardware may benefit from different floating-point precision (FP32) handling on the CPU or different library optimizations during the training phase.
2. **Robustness of the Proposed Pipeline:** While the CNN + CLAHE pipeline achieved its absolute peak of 97.84% on the local GPU, it remained highly effective on the DGX server with an accuracy of 96.16%. The variation is normal due to variations in randomized weight initialization and hardware-specific tensor management.

3. **Conclusion on Reliability:** Maintaining a consistency of >95% accuracy across entirely different processing architectures (Consumer GPU vs. Enterprise CPU) confirms the reliability of the MedVision-AI engine for diverse healthcare infrastructures.

5.6 Comparative Accuracy Visualization

To address the graphical benchmarking, the MedVision-AI dashboard features a comparative accuracy chart (see Figure 5.6). This visualization illustrates how our fine-tuned models—specifically the CNN+CLAHE pipeline—consistently outperform benchmarks reported in existing literature.




Figure 5.6 Performance Dashboard Benchmarks.

The detailed metrics and interpretation for each architecture are broken down in Figure 5.7.




Figure 5.7 Metrics Interpretation.

Beyond real-time analysis, the system maintains clinical accountability through a session history found in the Analysis Log (Figure 5.8). Finally, the system's end-to-end utility is demonstrated by the Automated Pathology Report (PDF) export, as illustrated in Figure 5.9.




Figure 5.8 Diagnosis Session History.

CLINICAL AI DIAGNOSIS & PATHOLOGY REPORT

Patient/Case Information

- **Analytical Method:** Neural Network Classification (ResNet50/VGG/CNN)
- **Primary Prediction:** Cancer
- **Statistical Confidence:** 88.52%
- **XAI Visualization:** Grad-CAM Activation Heatmap

Explainable AI (XAI) Findings

Based on the high-intensity activation centroids in the final convolutional layers:

The VGG19 architecture identified high-entropy activation patterns within the deeper convolutional blocks.

These clusters correlate with irregular structural density and architectural distortion characteristic of malignant lesions.

Clinical Context (Retrieved via RAG)

The following insights were retrieved from the medical knowledge base based on the suspicion of Cancer:

BI-RADS Assessment Categories

The Breast Imaging-Reporting and Data System (BI-RADS) is a standard for reporting:

- **BI-RADS 0:** Incomplete. Further imaging (e.g., spot compression, magnification, or ultrasound) is required.
- **BI-RADS 1:** Negative. Symmetrical breasts with no masses, architectural distortion, or suspicious calcifications.
- **BI-RADS 2:** Benign Findings. Includes secretory calcifications, simple cysts, and fat-containing lesions (hamartomas).
- **BI-RADS 3:** Probably Benign. <2% risk of malignancy. Recommend 6-month interval follow-up.
- **BI-RADS 4:** Suspicious. 2% to 95% likelihood of malignancy. Biopsy is typically indicated.
- **BI-RADS 5:** Highly Suggestive of Malignancy. >95% likelihood. Biopsy is mandatory.

DISCLAIMER: This report is generated by an AI assistant for research purposes only. It should be reviewed by a qualified medical professional before clinical action.

Page 1/2

Figure 5.9 PDF Pathology Report Sample.

5.7 Synthesis of Results

As evidenced by the benchmarks, our MedVision-AI framework consistently outperforms previous literature. The primary takeaway from these experiments is that while CNN architectures are powerful, their ultimate clinical utility is unlocked only when combined with rigorous preprocessing (CLAHE) and

cloud-optimized hyperparameter tuning. This multi-layered approach ensures that the system is not just an experimental prototype, but a viable diagnostic tool for real-world clinical implementation. As evidenced by the benchmarks, our MedVision-AI framework consistently outperforms previous literature. The primary takeaway from these experiments is that while CNN architectures are powerful, their ultimate clinical utility is unlocked only when combined with rigorous preprocessing (CLAHE) and cloud-optimized hyperparameter tuning. This multi-layered approach ensures that the system is not just an experimental prototype, but a viable diagnostic tool for real-world clinical implementation.



CHAPTER 6

CONCLUSION

6.1 Summary of Findings and The Solution

This research successfully developed a hybrid KNN-CNN framework, culminating in the MedVision-AI platform, specifically designed to address the critical challenges of automated breast cancer diagnosis. Through our experimental analysis, we confirmed that while both models are highly effective, the deep feature extraction capabilities of CNN architectures (VGG16 and VGG19) consistently outperform traditional KNN baselines—reaching a peak accuracy of 97.84% with CLAHE.

The core "Solution" presented in this thesis lies in the synergy between interpretability and predictive depth. By leveraging KNN's localized transparency alongside the high-precision abstraction of fine-tuned CNNs, we have provided a balanced diagnostic pipeline that is both medically reliable and technically scalable. The integration of this framework into a global cloud architecture (Google Antigravity) successfully bridges the historical gap between laboratory-based research and real-world clinical deployment. As established in the introductory chapters, this study fulfills its core objectives by delivering high-performance results that exceed literature benchmarks while prioritizing clinical trust through Grad-CAM explainability.

6.2 Limitations and Future Work

While the current framework provides a robust solution, several avenues for further enhancement remain. Future research will focus on:

1. Architectural Evolution: Incorporating next-generation architectures such as EfficientNet, ResNet-V2, and Vision Transformers (ViTs) to explore further improvements in classification precision and computational efficiency.
2. Adaptive Preprocessing: Exploring dynamic and adaptive contrast enhancement techniques that can automatically adjust to the specific pixel distribution of different imaging hardware.
3. Multimodal Integration: Enhancing the RAG (Retrieval-Augmented Generation) reporting layer to include multimodal clinical data, such as patient history and genomic markers, for a more holistic diagnostic profile.

In conclusion, the MedVision-AI framework proves that the integration of deep learning, cloud computing, and rigorous image preprocessing represents a significant advancement in the global fight against breast cancer.

6.3 Ethical Implications and Human-AI Collaboration

This research concludes with a critical reflection on the ethical deployment of AI in oncology. The goal of MedVision-AI is not to replace the radiologist, but to empower them with a "second pair of eyes."

1. Institutional Trust: By providing Grad-CAM heatmaps, we address the legal and moral requirement for transparency in medical decisions. The visual evidence allows for "Human-in-the-Loop" validation, ensuring that AI-driven insights are verified by human expertise before clinical action is taken.

2. Global Equity: The cloud-integrated nature of this framework ensures that high-quality cancer diagnostics are no longer restricted to wealthy urban centers. This study demonstrates that through the democratic use of cloud computing (Google Antigravity), we can provide equitable health outcomes for women in resource-limited regions, successfully fulfilling the broader mission of National Quemoy University to support global public health.



REFERENCES

- [1] World Health Organization (WHO), "Breast Cancer," 2024.
- [2] A. K. Barzan et al., "Mammogram Mastery: A Robust Dataset for Breast Cancer Detection," Mendeley Data, Apr. 2024.
- [3] Z. Zhu et al., "A Survey of Convolutional Neural Networks in Breast Cancer," CMES, vol. 136, 2022.
- [4] H. S. Das et al., "Breast Cancer Detection: CNNs vs. Shallow Networks," *Frontiers in Genetics*, 2023.
- [5] M. T. R. et al., "Optimized CNNs with ReduceLROnPlateau," IJCI Systems, 2024.
- [6] B. Abunasser et al., "Convolution Neural Network for breast cancer detection and classification using Deep learning," Asian Pacific Journal of Cancer Prevention, Feb. 2023.
- [7] S. Sharma et al., "Breast Cancer Detection Using Machine Learning Algorithms," 2018 International Conference on Computational Techniques, Electronics and Mechanical Systems (CTEMS), Dec. 2018.
- [8] M. A. Rufai et al., "MACHINE LEARNING MODEL FOR BREAST CANCER DETECTION," FUDMA JOURNAL OF SCIENCES, Feb. 2023.
- [9] T. Farjana, A. F. A. Momen, and F. Al-Amin, "Deep Learning Algorithm for Breast Masses Classification in Mammograms," ResearchGate, May 2020.
- [10] S. Saranya and S. Vijayarani, "Breast Cancer Classification Using Hybrid Features and Deep Neural Networks," *Research Square*, May 2023.
- [11] M. Mishra, M. H. Kolekar, and S. Sengupta, "Deep Learning for Breast Cancer Classification: Enhanced Tangent Function," *arXiv preprint arXiv:2108.04663*, 2021.
- [12] T. S. Kumar, G. Sridhar, D. Manju, P. Subhash, and G. Nagaraju, "Breast Cancer Classification Using ResNet50," Journal of Electrical Systems, 2023.
- [13] K. Simonyan and A. Zisserman, "Very Deep Convolutional Networks for Large-Scale Image Recognition," arXiv preprint arXiv:1409.1556, 2014. (Foundational VGG Paper)
- [14] K. He, X. Zhang, S. Ren, and J. Sun, "Deep Residual Learning for Image

Recognition," Proceedings of the IEEE conference on computer vision and pattern recognition, 2016. (Foundational ResNet Paper)

[15] R. R. Selvaraju et al., "Grad-CAM: Visual Explanations from Deep Networks via Gradient-based Localization," Proceedings of the IEEE international conference on computer vision, 2017. (Foundational Grad-CAM Paper)

[16] K. Zuiderveld, "Contrast Limited Adaptive Histogram Equalization," Graphics Gems IV, 1994. (Original CLAHE Source)

[17] J. Suckling et al., "The Mammographic Image Analysis Society Digital Mammogram Database," Excerpta Medica. International Congress Series, vol. 1069, 1994. (MIAS Dataset Origin)

[18] M. Heath et al., "The Digital Database for Screening Mammography," Proceedings of the 5th International Workshop on Digital Mammography, 2000. (DDSM Dataset Reference)

[19] M. Abadi et al., "TensorFlow: Large-Scale Machine Learning on Heterogeneous Systems," 2015.

[20] F. Chollet et al., "Keras," 2015. [Online]. Available: <https://keras.io>

[21] S. Ramírez, "FastAPI," 2018. [Online]. Available: <https://github.com/tiangolo/fastapi>

[22] W. Samek, T. Wiegand, and K. R. Müller, "Explainable Artificial Intelligence: Understanding, Visualizing and Interpreting Deep Learning Models," arXiv preprint arXiv:1708.08296, 2017.

[23] N. Srivastava et al., "Dropout: A Simple Way to Prevent Neural Networks from Overfitting," Journal of Machine Learning Research, vol. 15, no. 1, 2014.

[24] S. Ioffe and C. Szegedy, "Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift," International Conference on Machine Learning, 2015.

[25] D. P. Kingma and J. Ba, "Adam: A Method for Stochastic Optimization," arXiv preprint arXiv:1412.6980, 2014. (Standard AI Optimizer Reference)

APPENDICES

APPENDIX A

AI Model Core Implementation (Python).....	55
--	----

APPENDIX B

Backend API Services (FastAPI).....	58
-------------------------------------	----

APPENDIX C

Frontend Component Architecture (React/JSX).....	60
--	----

APPENDIX D


Data Preprocessing Utilities (Python/OpenCV)	62
--	----

APPENDIX E

Vector Similarity Search for RAG (FAISS/Sentence-Transformers).....	64
---	----

APPENDIX F

Tailwind CSS Interface Configuration	66
--	----



APPENDIX A

AI Model Core Implementation (Python)



Appendix A - AI Model Core Implementation (Python)

The following code snippet illustrates the Grad-CAM (Explainable AI) implementation within the MedVision-AI framework.

```
import tensorflow as tf
import numpy as np
import cv2

def make_gradcam_heatmap(img_array, model, last_conv_layer_name,
pred_index=None):
    # First, we create a model that maps the input image to the
activations
of the last conv layer as well as the output predictions
    grad_model = tf.keras.models.Model(
        [model.inputs],
        [model.get_layer(last_conv_layer_name).output, model.output]
    )

    # Then, we compute the gradient of the top predicted class for
our input image
with respect to the activations of the last conv layer
    with tf.GradientTape() as tape:
        last_conv_layer_output, preds = grad_model(img_array)
        if pred_index is None:
            pred_index = tf.argmax(preds[0])
        class_channel = preds[:, pred_index]

    # This is the gradient of the output neuron (top predicted or
chosen)
with regard to the output feature map of the last conv layer
    grads = tape.gradient(class_channel, last_conv_layer_output)

    # This is a vector where each entry is the mean intensity of
the gradient
over a specific feature map channel
    pooled_grads = tf.reduce_mean(grads, axis=(0, 1, 2))

    # We multiply each channel in the feature map array
by "how important this channel is" with regard to the top
predicted class
```

```
# then sum all the channels to obtain the heatmap class  
activation  
last_conv_layer_output = last_conv_layer_output[0]  
heatmap = last_conv_layer_output @ pooled_grads[..., tf.newaxis]  
heatmap = tf.squeeze(heatmap)  
  
# For visualization purpose, we will also normalize the heatmap  
between 0 & 1  
heatmap = tf.maximum(heatmap, 0) / tf.math.reduce_max(heatmap)  
return heatmap.numpy()
```



APPENDIX B

Backend API Services (FastAPI)



Appendix B - Backend API Services (FastAPI)

MedVision-AI utilizes a low-latency FastAPI backend to provide Inference-as-a-Service.

```
from fastapi import FastAPI, UploadFile, File
from pydantic import BaseModel
import tensorflow as tf

app = FastAPI(title="MedVision-AI Backend")

# Load pre-trained VGG16+CLAHE model
model = tf.keras.models.load_model("medvision_vgg16_clahe.h5")

@app.post("/api/v1/predict")
async def predict_mammogram(file: UploadFile = File(...)):
    # 1. Read and preprocess the image
    contents = await file.read()
    processed_image = preprocess_image(contents) # Custom
    # resize/CLAHE

    # 2. Dual-path inference (CNN)
    prediction = model.predict(processed_image)

    # 3. Generate XAI Evidence
    heatmap = make_gradcam_heatmap(processed_image, model,
    "block5_conv3")

    # 4. Trigger RAG Report
    clinical_report = rag_engine.generate(prediction)

    return {
        "diagnosis": "Malignant" if prediction > 0.5 else "Benign",
        "confidence": float(prediction),
        "xai_overlay_url": "/static/temp/heatmap.png",
        "report": clinical_report
    }
```

APPENDIX C

Frontend Component Architecture (React/JSX)



Appendix C - Frontend Component Architecture (React/JSX)

The clinician interface provides a visual dashboard for diagnostic validation.

```
import React, { useState } from 'react';

const DiagnosisDashboard = () => {
  const [image, setImage] = useState(null);
  const [result, setResult] = useState(null);

  const handleDiagnosis = async () => {
    const formData = new FormData();
    formData.append("file", image);

    const response = await fetch("/api/v1/predict", {
      method: "POST",
      body: formData
    });
    const data = await response.json();
    setResult(data);
  };

  return (
    <div className="p-10 glassmorphism-bg">
      <h1>MedVision-AI Portal</h1>
      <input type="file" onChange={(e) =>
setImage(e.target.files[0])} />
      <button onClick={handleDiagnosis}>Begin Automated
Diagnosis</button>

      {result && (
        <div className="flex gap-4">
          <img src={URL.createObjectURL(image)}
alt="Original" />
          <img src={result.xai_overlay_url} alt="XAI
Heatmap" />
          <p className="clinical-
report">{result.report}</p>
        </div>
      )}
    </div>
  );
};
```

APPENDIX D

Data Preprocessing Utilities (Python/OpenCV)



Appendix D - Data Preprocessing Utilities (Python/OpenCV)

The core image enhancement module for the MedVision-AI pipeline.

```
import cv2
import numpy as np

def apply_medvision_enhancement(input_image_path, target_size=(128,
128)):
    # 1. Load image in grayscale
    img = cv2.imread(input_image_path, cv2.IMAGING_GRAYSCALE)
    # 2. Resizing to model specification
    img_resized = cv2.resize(img, target_size)
    # 3. Applying CLAHE (Clip Limit: 2.0, Tile Grid: 8x8)
    clahe = cv2.createCLAHE(clipLimit=2.0, tileGridSize=(8, 8))
    enhanced_img = clahe.apply(img_resized)
    # 4. Normalization for CNN Input range [0, 1]
    normalized_img = enhanced_img.astype('float32') / 255.0
    # 5. Adding channel dimension for RGB-weighted models
    final_tensor = np.stack((normalized_img,)*3, axis=-1)
    return np.expand_dims(final_tensor, axis=0)
```



APPENDIX E

Vector Similarity Search for RAG (FAISS/Sentence-Transformers)



Appendix E - Vector Similarity Search for RAG (FAISS/Sentence-Transformers)

Implementation of the clinical knowledge retrieval layer.

```
from sentence_transformers import SentenceTransformer
import faiss

class ClinicalRAG:
    def __init__(self, knowledge_base_path):
        self.model = SentenceTransformer('all-MiniLM-L6-v2')
        self.index = faiss.read_index(knowledge_base_path)
        self.templates = ["Normal Mammogram...", "Suspicious
Malignancy...", "Ductal Carcinoma..."]

    def generate_grounded_report(self, pred_prob, confidence):
        # 1. Create search query based on prediction
        query = f"Mammogram finding with probability {pred_prob}
and {confidence}"
        query_vector = self.model.encode([query])
        # 2. Search FAISS index for WHO standards
        D, I = self.index.search(query_vector, k=1)
        # 3. Synthesize final clinical narrative
        return f"MedVision-AI Analysis: {self.templates[I[0][0]]}.
Confidence Score: {confidence*100}%"
```

NQU

APPENDIX F

Tailwind CSS Interface Configuration



Appendix F - Tailwind CSS Interface Configuration

Glassmorphism styling used in the React Frontend.

```
module.exports = {
  theme: {
    extend: {
      backdropFilter: {
        'none': 'none',
        'blur': 'blur(20px)',
      },
      colors: {
        'med-blue': '#0A192F',
        'med-teal': '#64FFDA',
        'glass-white': 'rgba(255, 255, 255, 0.05)',
      },
      boxShadow: {
        'glass-card': '0 8px 32px 0 rgba(31, 38, 135, 0.37)',
      },
    },
  },
}
```

The logo for NQU (National Qualifications University) is displayed. It features the letters 'NQU' in a bold, white, sans-serif font, centered within a light orange rounded rectangle. Above the text, there are three white, curved lines that suggest a stylized wave or a smile.

Authors' Background



Shi-Han Huang is currently pursuing a Master's degree in the Department of Computer Science and Information Engineering at National Quemoy University (NQU), Taiwan. He earned his Bachelor degree in Computer Science and Information Engineering from Ciputra University, Surabaya, Indonesia, in 2024. During his undergraduate studies, he completed an internship at the Apple Developer Academy @ UC, where he gained hands-on experience in innovative software development. His current research focuses on the application of Artificial Intelligence in Medical Imaging, with a particular interest in developing intelligent systems to support medical diagnostics.



Dr. Hsi-Chieh Lee is a Professor of the Department of Computer Science and Information Engineering at National Quemoy University (NQU) in Taiwan. He has Ph.D. degrees in the Department of Computer Science and the Department of Engineering Management from the University of Missouri-Rolla, USA. He also got an M.S. in Computer Science from the same institution and a B.S. in Mathematics from National Taiwan University. Dr. Lee has a wide range of research interests in artificial intelligence, biomedical informatics, renewable energy informatics, digital humanities, and high-performance computing. In addition, throughout his career, he has had many academic and administrative roles, which include time as Dean of the Academic Affairs, the College of Humanities and Arts, and the College of Health and Nursing at NQU. He also served as a professor and department chair at Yuan Ze University. In 2023, he was nominated as a full member of SIGMA XI, the Scientific Research Honor Society.